Causal Estimation for Text Data with (Apparent) Overlap Violations

Lin Gui[1] and Victor Veitch[1]

[1]*The University of Chicago*

Abstract

Consider the problem of estimating the causal effect of some attribute of a text document; for example: what effect does writing a polite vs. rude email have on response time? To estimate a causal effect from observational data, we need to adjust for confounding aspects of the text that affect both the treatment and outcome—e.g., the topic or writing level of the text. These confounding aspects are unknown a priori, so it seems natural to adjust for the entirety of the text (e.g., using a transformer). However, causal identification and estimation procedures rely on the assumption of overlap: for all levels of the adjustment variables, there is randomness leftover so that every unit could have (not) received treatment. Since the treatment here is itself an attribute of the text, it is perfectly determined, and overlap is apparently violated. The purpose of this paper is to show how to handle causal identification and estimation in the presence of apparent overlap violations. In brief, the idea is to use supervised representation learning to produce a data representation that preserves confounding information while eliminating information that is only predictive of the treatment. This representation then suffices for adjustment and satisfies overlap. Adapting results on non-parametric estimation, we show that this procedure yields a low-bias estimator that admits valid uncertainty quantification under weak conditions. Empirical results show reductions in bias and strong improvements in uncertainty quantification relative to the natural (transformer-based) baseline. Code, demo data and a tutorial are available at https://github.com/gl-ybnbxb/TI-estimator.

1 Introduction

We consider the problem of estimating the causal effect of an attribute of a passage of text on some downstream outcome. For example, what is the effect of writing a polite or rude email on the amount of time it takes to get a response? In principle, we might hope to answer such questions with a randomized experiment. However, this can be difficult in practice—e.g., if poor outcomes are costly or take a long to gather. Accordingly, in this paper, we will be interested in estimating such effects using observational data.

There are three steps to estimating causal effects using observational data. First, we need to specify a concrete causal quantity as our estimand. That is, give a formal quantity target of estimation corresponding to the high-level question of interest. The next step is causal identification: we need to prove that this causal estimator can, in principle, be estimated using only observational data. The standard approach for identification relies on adjusting for confounding variables that affect both the treatment and the outcome. For identification to hold, our adjustment variables must satisfy two conditions: unconfoundedness and overlap. The former requires the adjustment variables contain sufficient information on all common causes. The latter requires that the adjustment variable does not contain enough information about treatment assignment to let us perfectly predict it. Intuitively, to disentangle the effect of treatment from the effect of confounding, we must observe each treatment state at all levels of confounding. The final step is estimation using a finite data sample. Here, overlap also turns out to be critically important as a major determinant of the best possible accuracy (asymptotic variance) of the estimator

[Che+16].

Since the treatment is a linguistic property, it is reasonable to assume that text data has information about all common causes of the treatment and the outcome. In the email example, although there are many other factors, such as the recipient's schedule, affecting the response time not shown in the email, the information affecting both politeness and the response time should be included in the content of the email. Thus, it is easy to satisfy unconfoundedness in the text setting by adjusting for all the text as the confounding part. However, doing so brings about overlap violation. Since the treatment is a linguistic property determined by the text, the probability of treatment given any text is either 0 or 1. The polite/ rude tone is determined by the text itself. Therefore, overlap does not hold if we naively adjust for all the text as the confounding part. This problem is the main subject of this paper. Or, more precisely, our goal is to find a causal estimand, causal identification conditions, and estimation procedure that will allow us to effectively estimate causal effects even in the presence of such (apparent) overlap violations.

In fact, there is an obvious first approach: simply use a standard plug-in estimation procedure that relies only on modeling the outcome from the text and treatment variables. In particular, do not make any explicit use of the propensity score, the probability each unit is treated. Pryzant et al. [Pry+20] use an approach of this kind and show it is reasonable in some situations. Indeed, we will see in Section 3 and 4 that this procedure can be interpreted as a point estimator of a natural causal effect. However, there is a major drawback: it is unclear how to quantify the uncertainty of this kind of estimate. In principle, we could bootstrap. However, in practice, we would like to estimate the outcome from the text using large (transformer-based) language models. Refitting such models many times is prohibitively computationally expensive. We might also hope that we can simply ignore the model fitting uncertainty, and compute the uncertainty remaining in the estimator after conditioning on a particular model fit. Unfortunately, this is only valid under very strong conditions on how quickly the model learns the true text-outcome relationship. In practice, this procedure works very poorly, dramatically underestimating the uncertainty; see Section 5.

The contribution of this paper is a method for estimating causal effects in text that is both efficient and allows accurate uncertainty quantification under weak conditions without requiring model refitting. The main idea is to break estimation into a two-stage procedure, where in the first stage we learn a representation of the text that preserves enough information to account for confounding, but throws away enough information to avoid overlap issues. Then, we use this representation as the adjustment variables in a standard double machine-learning estimation procedure [Che+16], [Che+17a]. More precisely, the contributions of this paper are:

1. We give a formal causal estimand corresponding to the text-attribute question. We show this estimand is causally identified under weak conditions, even in the presence of overlap issues.
2. We show how to efficiently estimate this quantity using the adapted double-ML technique just described. We show that this estimator admits a central limit theorem (i.e., valid uncertainty quantification) under weak conditions on the rate at which the ML model learns the text-outcome relationship.
3. We test the performance of this procedure empirically, finding significant improvements in bias and uncertainty quantification relative to the outcome-model-only baseline.

Related work The most related literature is on causal inference with text variables. Papers include treating text as treatment [Pry+20; WDSD18; Ega+18; FG16; WC19; TLP14]), as outcome [Ega+18; SG19], as confounder [VSB19; RSN20; Moz+20; KJO20], and discovering or predicting causality from text [PMB16; Tab+18; Bal+19; MC00]. There are also numerous applications using text to adjust for confounding: [OVK17; Hal17; KCG18; Sri+18; SG19; Sah+19; KF19; ZMDNM20]. Of these, Pryzant et al. [Pry+20] also address non-parametric estimation of the

causal effect of text attributes. We follow their formalization of the problem, extending it to handle overlap violations. Their focus is primarily on misestimation of the treatments. They do not address uncertainty quantification, which is our main motivation.

This paper also relates to work on causal estimation with (near) overlap violations. D'Amour et al. [D'A+21] points out high-dimensional adjustment [e.g., Ras+11; Lou+17; Li+16; Ath+17] suffers from overlap issues. Extra assumptions such as sparsity and trimming are often needed to meet the overlap condition. These results do not directly apply here because we assume there exists a low-dimensional summary that suffices to handle confounding.

D'Amour and Franks [DF21] studies summary statistics that suffice for identification, which they call deconfounding scores. The supervised representation learning approach in this paper can be viewed as an extremal case of the deconfounding score. However, the motivation and technical results are distinct. That work focuses on finding all deconfounding scores in a linear-gaussian setting. Here, our concern is with uncertainty quantification in the presence of overlap violations when learning with large neural models. Hence, the papers are complimentary.

2 Notation and Problem Setup

We follow the causal setup of Pryzant et al. [Pry+20]. We are interested in estimating the causal effect of treatment A on outcome Y. For example, how does writing a negative sentiment (A) review (X) affect product sales (Y)? There are two immediate challenges to estimating such effects with observed text data. First, we do not actually observe A, which is the intent of the writer. Instead, we only observe $\tilde{A}$, a version of A that is inferred from the text itself. In this paper, we will assume that $A = \tilde{A}$ almost surely—e.g., a reader can always tell if a review was meant to be negative or positive. This assumption is often reasonable, and follows [Pry+20]. The next challenge is that the treatment may be correlated with other aspects of the text (Z) that are also relevant to the outcome—e.g., the product category of the item being reviewed. Such Z can act as confounding variables, and must somehow be adjusted for in a causal estimation problem.

Specifically, a writer writes a text X based on some linguistic properties A and Z. Note that A and Z are possibly correlated with each other. $\tilde{A}$ is the treatment perceived by a reader. In the product example, a customer writes a review X of a product based on his/her satisfaction A and some other properties Z, such as product types. $\tilde{A}$ is the attitude another customer perceives from the review. It is reasonable to assume that the writer is able to express himself/herself clearly, so the linguistic property A he/she uses can be successfully received by the reader. Mathematically, the assumption is that $A = \tilde{A}$ almost surely. The content of this text document X and the perceived treatment $\tilde{A}$ jointly determines the outcome Y, namely the customer's purchase behavior (buy or not).

Each unit (A_i, Z_i, X_i, Y_i) is independently and identically sampled from an unknown distribution P. Figure 1 shows the causal relationships among variables, where solid arrows represent causal relations, and the dotted line represents possible correlations between two variables. We assume that text X contains all common causes of $\tilde{A}$ and the outcome Y.

3 Identification and Causal estimand

The first task is to translate the qualitative causal question of interest—what is the effect of A on Y—into a causal estimand. This estimand must both be faithful to the qualitative question and be identifiable from observational data under reasonable assumptions. The key challenges here are that we only observe $\tilde{A}$ (not A itself), there are unknown confounding variables influencing the text, and $\tilde{A}$ is a deterministic function of the text, leading to overlap violations if we naively adjust for all the text. Our high-level idea is to split the text into abstract (unknown) parts

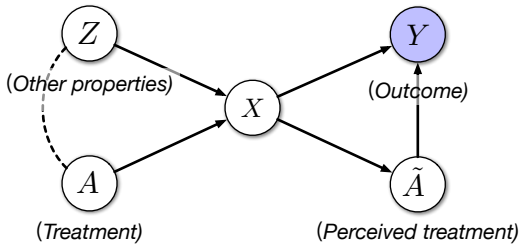

Figure 1: The causal DAG of the problem. A writer writes a text document X based on linguistic properties A and Z, where A is the treatment in the causal problem. A and Z cannot be observed directly in data and can only be seen via text. The dotted line represents possible correlation between A and Z. A reader perceives the treatment $\tilde{A}$ from the text. The perceived treatment $\tilde{A}$ together with contents of X determine the outcome Y.

depending on whether they are confounding—affect both $\tilde{A}$ and Y—or whether they affect $\tilde{A}$ alone. The part of the text that affects only $\tilde{A}$ is not necessary for causal adjustment, and can be thrown away. If this part contains "enough" information about $\tilde{A}$, then throwing it away will eliminate our ability to perfectly predict $\tilde{A}$, thus fixing the overlap issue. We now turn to formalizing this idea, showing how it can be used to define an estimand and to identify this estimand from observational data.

Causal model The first idea is to decompose the text into three parts: one part affected by only A, one part affected interactively by A and Z, and another part affected only by Z. We use X_A, $X_{A\wedge Z}$ and X_Z to denote them, respectively; see Figure 2 for the corresponding causal model. Note that there could be additional information in the text in addition to these three parts. However, since they are irrelevant to both A and Z, we do not need to consider them in the model.

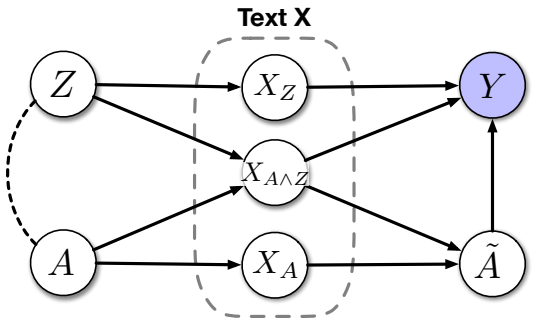

Figure 2: A more sophisticated causal model with the decomposition of text X. X_A, $X_{A\wedge Z}$, and X_Z are parts of the text affected by only A, both A and Z, and only Z, respectively. A and Z are linguistic properties a writer based on and thus cannot be observed directly from data. When investigating the causal relationship between $\tilde{A}$ and Y, $(X_{A\wedge Z}, X_Z)$ is a confounding part satisfying both unconfoundedess and overlap.

NDE The treatment A affects the outcome through two paths. Both "directly" through X_A—the part of the text determined just by the treatment—and also through a path going through $X_{A\wedge Z}$—the part of the text that relies on interaction effects with other factors. Our formal causal effect aims at capturing the effect of A through only the first, direct, path.

$$\text{NDE} := \mathbb{E}_{X_{A\wedge Z}, X_Z | A=1}\left[\mathbb{E}[Y \mid X_{A\wedge Z}, X_Z, \text{do}(A=1)] - \mathbb{E}[Y \mid X_{A\wedge Z}, X_Z, \text{do}(A=0)]\right]. \tag{3.1}$$

Here, do is Pearl's do notation, and the estimand is a variant of the natural direct effect [Pea09]. Intuitively, it can be interpreted as the expected change in the outcome induced by changing the treatment from 1 to 0 while keeping part of the text affected by Z the same as it would have been had we set $A = 1$. This is a reasonable formalization of the qualitative "effect of A on Y". Of course, it is not the only possible formalization. Its advantage is that, as we will see, it can be identified and estimated under reasonable conditions.

Identification To identify NDE we must rewrite the expression in terms of observable quantities. There are three challenges: we need to get rid of the do operator, we don't observe A (only $\tilde{A}$), and the variables $X_{A\wedge Z}, X_Z$ are unknown (they are latent parts of X).

Informally, the identification argument is as follows. First, $X_{A\wedge Z}, X_Z$ block all backdoor paths (common causes) in Figure 2. Moreover, because we have thrown away X_A, we now satisfy overlap. Accordingly, the do operator can be replaced by conditioning following the usual causal-adjustment argument. Next, $A = \tilde{A}$ almost surely, so we can just replace A with $\tilde{A}$. Now, our estimand has been reduced to:

$$\tilde{\text{NDE}} := \mathbb{E}_{X_{A\wedge Z}, X_Z|\tilde{A}=1}\Big[\mathbb{E}[Y \mid X_{A\wedge Z}, X_Z, \tilde{A}=1] - \mathbb{E}[Y \mid X_{A\wedge Z}, X_Z, \tilde{A}=0)]\Big]. \tag{3.2}$$

The final step is to deal with the unknown $X_{A\wedge Z}, X_Z$. To fix this issue, we first define the *conditional outcome* Q according to:

$$Q(\tilde{A}, X) := \mathbb{E}(Y \mid \tilde{A}, X). \tag{3.3}$$

A key insight here is that, subject to the causal model in Figure 2, we have $Q(\tilde{A}, X) = \mathbb{E}(Y \mid \tilde{A}, X_{A\wedge Z}, X_Z)$. But this is exactly the quantity in (3.2). Moreover, $Q(\tilde{A}, X)$ is an observable data quantity (it depends only on the distribution of the observed quantities). In summary:

Theorem 1. *Assume the following:*

1. *(Causal structure) The causal relationships among A, $\tilde{A}$, Z, Y, and X satisfy the causal DAG in Figure 2*
2. *(Overlap)* $0 < \mathrm{P}(A = 1 | X_{A\wedge Z}, X_Z) < 1$
3. *(Intention equals perception)* $A = \tilde{A}$ *almost surely.*

Then, the NDE is identified from observational data as

$$\text{NDE} = \tau^{\text{NDE}} := \mathbb{E}_{X|\tilde{A}=1}\Big[\mathbb{E}[Y \mid \eta(X), \tilde{A}=1] - \mathbb{E}[Y \mid \eta(X), \tilde{A}=0]\Big], \tag{3.4}$$

where $\eta(X) := (f(Q(0,X)), h(Q(1,X)))$*, and* $f, h : \mathbb{R} \to \mathbb{R}$ *are any invertible functions on* $\mathbb{R}$.

We give the result in terms of an abstract sufficient statistic $\eta(X)$ to emphasize that the actual conditional expectation model is not required, only some statistic that is informationally equivalent.

4 Method

Our ultimate goal is to draw a conclusion about whether the treatment has a causal effect on the outcome. Following previous section, we have reduced this problem to estimating τ^{NDE}, defined in Theorem 1. The task now is to develop an estimation procedure, including uncertainty quantification.

4.1 Outcome only estimator

We start by introducing the naive outcome only estimator as a first approach to NDE estimation. The estimator is adapted from Pryzant et al. [Pry+20]. The observation here is that, taking $\eta(X) = (Q(0,X), Q(1,X))$ in (3.4), we have

$$\tau^{\text{NDE}} = \mathbb{E}_{X|A=1}[\ \mathbb{E}(Y \mid A=1, X) - \mathbb{E}(Y \mid A=0, X)\ ]. \tag{4.1}$$

Since $Q(A,X)$ is a function of the whole text data X, it is estimable from observational data. Namely, it is the solution to the square error risk:

$$Q = \operatorname*{argmin}_{\tilde{Q}} \mathbb{E}[(Y - \tilde{Q}(A,X)^2]. \tag{4.2}$$

With a finite sample, we can estimate Q as $\hat{Q}$ by fitting a machine-learning model to minimize the (possibly regularized) square error empirical risk. That is, fit a model using mean square error as the objective function. Then, a straightforward estimator is:

$$\hat{\tau}^{\mathrm{Q}} := \frac{1}{n_1} \sum_{i:A_i=1} \hat{Q}_1(X_i) - \hat{Q}_0(X_i), \tag{4.3}$$

where n_1 is the number of treated units.

As discussed in the introduction, this estimator is reasonable for point estimates but does not offer a simple approach for uncertainty quantification. A natural guess for an estimate of its variance is:

$$\hat{\mathrm{var}}(\hat{\tau}^{\mathrm{Q}}) := \frac{1}{n} \hat{\mathrm{var}}(\hat{Q}_1(X_i) - \hat{Q}_0(X_i) \mid \hat{Q}). \tag{4.4}$$

That is, just compute the variance of the mean conditional on the fitted model. This procedure yields asymptotically valid confidence intervals if the outcome model converges extremely quickly; i.e., if $\mathbb{E}[(\hat{Q} - Q)^2]^{\frac{1}{2}} = o(n^{-\frac{1}{2}})$. We could instead bootstrap, refitting $\hat{Q}$ on each bootstrap sample. However, with modern language models, this can be prohibitively computationally expensive.

4.2 Treatment Ignorant Effect Estimation (TI-estimator)

Following Theorem 1, it suffices to adjust for $\eta(X) = (Q(0,X), Q(1,X))$. Accordingly, we use the following pipeline. We first estimate $\hat{Q}_0(X)$ and $\hat{Q}_1(X)$ (using a neural language model), as with the outcome-only estimator. Then, we take $\hat{\eta}(X) := (\hat{Q}_0(X), \hat{Q}_1(X))$ and estimate $\hat{g}_\eta \approx \mathrm{P}(A = 1 \mid \hat{\eta})$. That is, we estimate the propensity score corresponding to the estimated representation. Finally, we plug the estimated $\hat{Q}$ and $\hat{g}_\eta$ into a standard double machine learning estimator [Che+16].

We describe the three steps in detail:

Q-Net In the first stage, we intend to estimate the conditional outcomes and hence obtain the estimated two-dimensional confounding vector $\hat{\eta}(X)$. We may use any approach for estimating the conditional outcomes. For concreteness, we will use the dragonnet architecture of Shi et al. [SBV19]. Specifically, we train DistilBERT [San+19] modified to include three heads, as shown in Section 4.2. Two of the heads correspond to $\hat{Q}_0(X)$ and $\hat{Q}_1(X)$ respectively. The final head is a single linear layer predicting the treatment (which can achieve perfect accuracy). The output of this head is not used for the estimation; its purpose is to force the DistilBERT representation to preserve all confounding information. This has been shown to improve causal estimation [SBV19; VSB19].

We train the model by minimizing the objective function

$$\mathscr{L}(\theta; \mathbf{X}) = \frac{1}{n} \sum_i \Big[\big(\hat{Q}_{a_i}(x_i; \theta) - y_i \big)^2 + \alpha \mathrm{CrossEntropy}\,(a_i, g_u(x_i)) + \beta \mathscr{L}_{\mathrm{mlm}}(x_i) \Big], \tag{4.5}$$

where θ are the model parameters, α, β are hyperparameters and $\mathscr{L}_{\mathrm{mlm}}(\cdot)$ is the masked language modeling objective of DistilBERT.

There is a final nuance. In practice, we split the data into K-folds. For each fold j, we train a model $\hat{Q}_{-j}$ on the other $K-1$ folds. Then, we make predictions for the data points in fold j using $\hat{Q}_{-j}$. Slightly abusing notation, we use $\hat{Q}_a(x)$ to denote the predictions obtained in this manner.

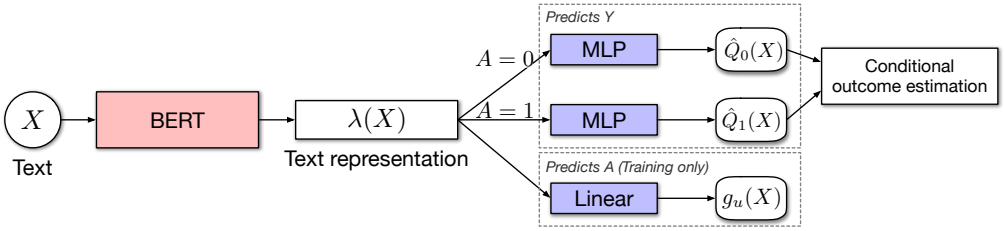

Figure 3: The architecture of Q-Net follows the dragonnet [SBV19] for estimation of $\hat{Q}$. Specifically, given representations $\lambda(X)$ from input text data, the Q-Net predicts Y for samples with $A=0$ and $A=1$ using two separate heads. A third head predicting A is also included for training, though the predictions are not used for estimation. Parameters in BERT and three prediction heads are trained together in an end-to-end manner.

Propensity score estimation Next, we define $\hat{\eta}(x) \coloneqq (\hat{Q}_0(x), \hat{Q}_1(x))$ and estimate the propensity score $\hat{g}_\eta(x) \approx \mathrm{P}(A=1 \mid \hat{\eta}(x))$. To do this, we fit a non-parametric estimator to the binary classification task of predicting A from $\hat{\eta}(X)$. The important insight here is that since $\hat{\eta}(X)$ is 2-dimensional, non-parametric estimation is possible at a fast rate. In Section 5, we try several methods and find that kernel regression usually works well.

We also define $g_\eta(X) \coloneqq \mathrm{P}(A=1 \mid \eta(X))$ as the idealized propensity score. The idea is that as $\hat{\eta} \to \eta$, we will also have $\hat{g}_\eta \to g_\eta$ so long as we have a valid non-parametric estimate.

NDE estimation The final stage is to combine the estimated outcome model and propensity score into a NDE estimator. To that end, we define the influence curve of τ^{NDE} as follows:

$$\phi(X; Q, g_\eta, \tau^{\mathrm{NDE}}) \coloneqq A \cdot (Y - Q(0,X)) - \frac{g_\eta(X)}{1-g_\eta(X)} \cdot (1-A) \cdot (Y - Q(0,X)) - \tau^{\mathrm{NDE}} \tag{4.6}$$

Then, the standard double machine learning estimator of τ^{NDE} [Che+16], and the confidence interval of this estimator, is given by

$$\hat{\tau}^{\mathrm{TI}} = \underset{1 \le i \le n}{\mathrm{mean}}\left(\hat{\phi}_i\right),\ CI^{\mathrm{TI}} = \left(\hat{\tau}^{\mathrm{TI}} - z_{1-\alpha/2}\hat{sd}(\hat{\phi}_i),\ \hat{\tau}^{\mathrm{TI}} + z_{1-\alpha/2}\hat{sd}(\hat{\phi}_i)\right), \tag{4.7}$$

where

$$\hat{\phi}_i = A_i \cdot \left(Y_i - \hat{Q}_0(X_i)\right) - \frac{\hat{g}_\eta(X_i)}{1-\hat{g}_\eta(X_i)} \cdot (1-A_i) \cdot \left(Y_i - \hat{Q}_0(X_i)\right),\ \ i = 1, \cdots, n, \tag{4.8}$$

$z_{1-\alpha/2}$ is the $\alpha/2$-upper quantile of the standard normal, and $\hat{sd}(\cdot)$ is the sample standard deviation.

Validity We now have an estimation procedure. It remains to given conditions under which this procedure is valid. In particular, we require that it should yield a consistent estimate and asymptotically correct confidence intervals.

Theorem 2. *Assume the following.*

1. *The mis-estimation of conditional outcomes can be bounded as follows*

$$\max_{a \in \{0,1\}} \mathbb{E}[(\hat{Q}_a(X) - Q(a,X))^2]^{\frac{1}{2}} = o(n^{-\frac{1}{4}}). \tag{4.9}$$

2. *The propensity score function $P(A=1|\cdot,\cdot)$ is Lipschitz continuous on $\mathbb{R}^2$, and $\exists\ \varepsilon > 0$, $\mathrm{P}\left(\varepsilon \le g_\eta(X) \le 1-\varepsilon\right) = 1$*
3. *The propensity score estimate converges at least as quickly as k nearest neighbor; i.e., $\mathbb{E}[\left(\hat{g}_\eta(X) - \mathrm{P}(A=1 \mid \hat{\eta}(X)\right)^2 \mid X]^{\frac{1}{2}} = O(n^{-\frac{1}{4}})$ [Gyö+02];*

4. *There exist positive constants C_1, C_2, c, and $q > 2$ such that*

$$\begin{aligned}
&\mathbb{E}[|Y|^q]^{\frac{1}{q}} \leq C_2, \\
&\sup_{\eta \in supp(\eta(X))} \mathbb{E}[(Y - Q(A,X)^2 \mid \eta(X) = \eta)] \leq C_2, \\
&\mathbb{E}[(Y - Q(A,X)^2)]^{\frac{1}{2}} \geq c, \\
&\max_{a \in \{0,1\}} \mathbb{E}[\left|\hat{Q}_a(X) - Q(a,X)\right|]^{\frac{1}{q}} \leq C_1.
\end{aligned}$$

Then, the estimator $\hat{\tau}^{TI}$ is consistent and

$$\sqrt{n}(\hat{\tau}^{TI} - \tau^{\text{NDE}}) \xrightarrow{d} \mathbb{N}(0, \sigma^2) \tag{4.10}$$

where $\sigma^2 = E\left(\phi(X; Q, g_\eta, \tau^{\text{NDE}})\right)^2$.

The proof is provided in Appendix A.

The key point from this theorem is that asymptotic normality requires only a weak condition on the convergence rate of Q. Intuitively, the reason is simply that, because $\hat{\eta}(X)$ is only 2-dimensional, it is always possible to nonparametrically estimate the propensity score from $\hat{\eta}$ at a fast rate—even naive KNN works! Effectively, this means the rate at which we estimate the true propensity score $g_\eta(X) = \mathrm{P}(A = 1 \mid \eta(X))$ is dominated by the rate at which we estimate $\eta(X)$, which is in turn determined by the rate for $\hat{Q}$. Now, the key property of the double ML estimator is that convergence only depends on the *product* of the convergence rates of $\hat{Q}$ and $\hat{g}$. Accordingly, we only need to estimate $\hat{Q}$ at the square root of the rate we needed for the naive Q-only procedure. This is much more plausible in practice. As we will see in Section 5, the TI-estimator dramatically improves the quality of the estimated confidence intervals.

Remark 3. In addition to the uncertainty quantification, there are some other (less important) advantages this estimation procedure inherits from the double ML estimator. First, if $\hat{Q}$ is consistent, then the estimator is nonparametrically efficient in the sense that no other nonparametric estimator has a smaller asymptotic variance. That is, the procedure using the data as efficiently as possible. Second, the estimator is still consistent even if $\hat{Q}$ is only consistent up to some invertible transformation. This is because $\hat{\eta}$ remains consistent in this case and the AIPTW is double robust. This property may be useful when $\hat{Q}$ is miscalibrated.

5 Experiments

We empirically study the method's capability to provide accurate causal estimates with good uncertainty quantification. Testing using semi-synthetic data (where ground truth causal effects are known), we find that the estimation procedure yields accurate causal estimates and confidence intervals. In particular, the TI-estimator has significantly lower bias and vastly better uncertainty quantification than the Q-only method.

Additionally, we study the effect of the choice of nonparametric propensity score estimator and the choice of double machine-learning estimator. These results are reported in Appendix C. Although these choices do not matter asymptotically, we find they have a significant impact in actual finite sample estimation. We find that, in general, kernel regression works well for propensity score estimation and the vanilla AIPTW corresponding to the NDE works well.

Finally, we reproduce the real-data analysis from [Pry+20]. We find that, with the more accurate uncertainty quantification, we are unable to reject the null hypothesis of no effect.

Table 1: The TI-estimator significantly improves both bias and coverage relative to the baseline. Tables show average bias and confidence interval coverage of NDE estimates, over 50 resimulations. The TI estimator $\hat{\tau}^{\text{TI}}$ displays higher accuracy/smaller bias of point estimate and much larger coverage proportions compared to outcome-only estimator $\hat{\tau}^{Q}$. The treatment level equals true NDE, which takes 1.0 (with causal effect) and 0.0 (without causal effect). Low and high noise level corresponds to γ set to 1.0 and 4.0. Low and high confounding level corresponds to β_c set to 50.0 and 100.0.

(a) Average bias

Noise:	Low				High			
True NDE:	1.0		0.0		1.0		0.0	
Confounding:	Low	High	Low	High	Low	High	Low	High
$\hat{\tau}^{Q}$	0.093	0.367	0.074	0.327	0.548	0.524	0.492	0.494
$\hat{\tau}^{\text{TI}}$	0.065	0.052	0.104	0.070	0.077	0.044	-0.038	0.084

(b) Coverage proportions of 95% confidence intervals

Noise:	Low				High			
True NDE:	1.0		0.0		1.0		0.0	
Confounding:	Low	High	Low	High	Low	High	Low	High
$\hat{\tau}^{Q}$	0%	0%	2%	0%	0%	0%	0%	0%
$\hat{\tau}^{\text{TI}}$	60%	86%	60%	86%	86%	88%	80%	84%

5.1 Amazon Reviews

Dataset We closely follow the setup of Pryzant et al. [Pry+20]. We use publicly available Amazon reviews for music products as the basis for our semi-synthetic data. We include reviews for mp3, CD and vinyl, and among these exclude reviews for products costing more than $100 or shorter than 5 words. The treatment A is whether the review is five stars ($A = 1$) or one/two stars ($A = 0$).

To have a ground truth causal effect, we must now simulate the outcome. To produce a realistic dataset, we choose a real variable as the confounder. Namely, the confounder C is whether the product is a CD ($C = 1$) or not ($C = 0$). Then, outcome Y is generated according to $Y \leftarrow \beta_a A + \beta_c(\pi(C) - \beta_o) + \gamma N(0,1)$. The true causal effect is controlled by β_a. We choose $\beta_a = 0.0,\ 1.0$ to generate data with and without causal effects. In this setting, β_t is the oracle value of our causal estimand. The strength of confounding is controlled by β_c. We choose $\beta_c = 50.0,\ 100.0$. The ground-truth propensity score is $\pi(C) = P(A = 1|C)$. We set it to have the value $\pi(0) = 0.8$ and $\pi(1) = 0.6$ (by subsampling the data). Finally, the noise level is controlled by γ; we choose 1.0 and 4.0 to simulate data with small and large noise. The final dataset has 10,685 data entries.

Protocol For the language model, we use pretrained `distilbert-base-uncased` model provided by the `transformers` package. The model is trained in the k-folding fashion with 5 folds. We apply the Adam optimizer [KB14] with a learning rate of $2e^{-5}$ and a batch size of 64. The maximum number of epochs is set as 20, with early stopped based on validation loss with a patience of 6. Each experiment is replicated with five different seeds and the final $\hat{Q}(a, x_i)$ predictions are obtained by averaging the predictions from the 5 resulting models. The propensity model is implemented by running the Gaussian process regression using `GaussianProcessClassifier` in the `sklearn` package with `DotProduct + WhiteKernel` kernel. (We choose different random state for the GPR to guarantee the convergence of the GPR.) The coverage experiment uses 50 replicates.

Results The main question here is the efficacy of the estimation procedure. Table 1 compares the outcome-only estimator $\hat{\tau}^{Q}$'s and the estimator $\hat{\tau}^{\text{TI}}$. First, the bias of the new method is significantly lower than the bias of the outcome-only estimator. This is particularly true where there is moderate to high levels of confounding. This improvement could be due either to the improved efficiency of the procedure, or its robustness against misestimation of Q. Next, we check actual coverage rates over 50 replicates of the experiment. First, we find that the naive approach

Table 2: Politeness does not has a causal effect on response time. Table displays different NDE estimates and their 95% confidence intervals. The unadjusted one is the difference of sample means of treatment (polite) group and control group. Only confidence interval of $\hat{\tau}_{\text{AIPTW}}$ covers 0.

Estimator	NDE	Confidence Interval
unadjusted $\hat{\tau}^{naive}$	-0.064	[-0.0983, -0.0304]
$\hat{\tau}^{Q}$	-0.0422	[-0.0441, -0.0403]
$\hat{\tau}^{TI}$	-0.025	[-0.058, 0.008]

for the outcome-only estimator fails completely. The nominal confidence interval almost never actually includes the true effect. It is wildly optimistic. By contrast, the confidence intervals from the new method often cover the true value. This is an enormous improvement over the baseline. Nevertheless, they still do not actually achieve their nominal (95%) coverage. This may be because the $\hat{Q}$ estimate is still not good enough for the asymptotics to kick in, and we are not yet justified in ignoring the uncertainty from model fitting. It is particularly striking that coverage is worse when there's confounding—this may reflect the fact that more confounding means there's more true "signal" in the data for the model to learn. (With no confounding, the outcome is unrelated to the text).

5.2 Application: Consumer Complaints to the Financial Protection Bureau

We follow the same pipeline of the real data experiment in [Pry+20, §6.2]. The dataset is consumers complaints made to the financial protection. Treatment A is politeness (measured using [YKT18]) and the outcome Y is a binary indicator of whether complaints receive a response within 15 days.

We use the same training procedure as for the simulation data, except we use only 3 folds. Table 2 shows point estimates and their 95% confidence intervals. Notice that both the naive and outcome-only estimators show a significant *negative* effect of politeness on response time. On the other hand, the more accurate AIPTW method has a confidence interval that covers 0, so we cannot reject the null hypothesis that consumers' politeness does not affect response time. This matches our intuitions that being more polite should not reduce the probability of receiving a timely reply. The naive and outcome-only estimates presumably support the (likely wrong) conclusion due to failing to sufficiently adjust for confounding, and failing to adequately quantify uncertainty.

6 Discussion

In this paper, we address the estimation of the causal effect of a text document attribute using observational data. The key challenge is that we must adjust for the text—to handle confounding—but adjusting for all of the text violates overlap. We saw that this issue could be effectively circumvented with a suitable choice of estimand and estimation procedure. In particular, we can find an estimand that corresponds to the qualitative causal question, and an estimator for this quantity that is both low bias and that admits valid confidence intervals under weak conditions. The procedure also circumvents the need for bootstrapping, which is prohibitively expensive in our setting.

There are some limitations. The actual coverage proportion of our estimator is below the nominal level. This is presumably due to the imperfect fit of the conditional outcome model. It would be interesting to try the procedure with larger-scale language models, or to find techniques to account for the uncertainty remaining in the model fit. For example, it may be possible to use a small number of model refits to quantify uncertainty induced by the choice of random seed.

Although we have focused on text in this paper, the problem of causal estimation with apparent overlap violation exists in any problem where we must adjust for unstructured and high-dimensional covariates. Another interesting direction for future work is to understand how analogous procedures work outside the text setting.

References

[Ath+17] S. Athey, G. Imbens, T. Pham, and S. Wager. "Estimating average treatment effects: supplementary analyses and remaining challenges". In: *American Economic Review* 5 (2017).

[Bal+19] A. Balashankar, S. Chakraborty, S. Fraiberger, and L. Subramanian. "Identifying predictive causal factors from news streams". In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*. 2019.

[Che+17a] V. Chernozhukov, D. Chetverikov, M. Demirer, E. Duflo, C. Hansen, and W. Newey. "Double/debiased/neyman machine learning of treatment effects". In: *American Economic Review* 5 (2017).

[Che+16] V. Chernozhukov, D. Chetverikov, M. Demirer, E. Duflo, C. Hansen, W. Newey, and J. Robins. "Double/debiased machine learning for treatment and causal parameters". In: *arXiv preprint arXiv:1608.00060* (2016).

[Che+17b] V. Chernozhukov, D. Chetverikov, M. Demirer, E. Duflo, C. Hansen, W. Newey, and J. Robins. "Double/debiased machine learning for treatment and structural parameters". In: *The Econometrics Journal* (2017).

[D'A+21] A. D'Amour, P. Ding, A. Feller, L. Lei, and J. Sekhon. "Overlap in observational studies with high-dimensional covariates". In: *Journal of Econometrics* 2 (2021).

[DF21] A. D'Amour and A. Franks. "Deconfounding scores: feature representations for causal effect estimation with weak overlap". In: *arXiv preprint arXiv:2104.05762* (2021).

[Ega+18] N. Egami, C. J. Fong, J. Grimmer, M. E. Roberts, and B. M. Stewart. "How to make causal inferences using texts". In: *arXiv preprint arXiv:1802.02163* (2018).

[FG16] C. Fong and J. Grimmer. "Discovery of treatments from text corpora". In: *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 2016.

[Gyö+02] L. Györfi, M. Kohler, A. Krzyzak, H. Walk, et al. *A distribution-free theory of nonparametric regression*. 2002.

[Hal17] A. Hall. "How hiring language reinforces pink collar jobs". In: *Textio Word Nerd* (2017).

[KF19] D. Karell and M. Freedman. "Rhetorics of radicalism". In: *American Sociological Review* 4 (2019).

[KJO20] K. A. Keith, D. Jensen, and B. O'Connor. "Text and causal inference: a review of using text to remove confounding from causal estimates". In: *arXiv preprint arXiv:2005.00649* (2020).

[KCG18] E. Kiciman, S. Counts, and M. Gasser. "Using longitudinal social media analysis to understand the effects of early college alcohol use". In: *Twelfth International AAAI Conference on Web and Social Media*. 2018.

[KB14] D. P. Kingma and J. Ba. "Adam: a method for stochastic optimization". In: *arXiv preprint arXiv:1412.6980* (2014).

[Li+16] S. Li, N. Vlassis, J. Kawale, and Y. Fu. "Matching via dimensionality reduction for estimation of treatment effects in digital marketing campaigns." In: *IJCAI*. 2016.

[Lou+17] C. Louizos, U. Shalit, J. M. Mooij, D. Sontag, R. Zemel, and M. Welling. "Causal effect inference with deep latent-variable models". In: *Advances in neural information processing systems* (2017).

[MC00] S. Mani and G. F. Cooper. “Causal discovery from medical textual data.” In: *Proceedings of the AMIA Symposium*. American Medical Informatics Association. 2000.

[Moz+20] R. Mozer, L. Miratrix, A. R. Kaufman, and L. J. Anastasopoulos. “Matching with text data: an experimental evaluation of methods for matching documents and of measuring match quality”. In: *Political Analysis* 4 (2020).

[OVK17] A. Olteanu, O. Varol, and E. Kiciman. “Distilling the outcomes of personal experiences: a propensity-scored analysis of social media”. In: *Proceedings of the 2017 ACM Conference on Computer Supported Cooperative Work and Social Computing*. 2017.

[Pea09] J. Pearl. *Causality*. 2nd. 2009.

[PMB16] F. M. del Prado Martin and C. Brendel. “Case and cause in icelandic: reconstructing causal networks of cascaded language changes”. In: *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 2016.

[Pry+20] R. Pryzant, D. Card, D. Jurafsky, V. Veitch, and D. Sridhar. “Causal effects of linguistic properties”. In: *arXiv preprint arXiv:2010.12919* (2020).

[Ras+11] J. A. Rassen, R. J. Glynn, M. A. Brookhart, and S. Schneeweiss. “Covariate selection in high-dimensional propensity score analyses of treatment effects in small samples”. In: *American journal of epidemiology* 12 (2011).

[RSN20] M. E. Roberts, B. M. Stewart, and R. A. Nielsen. “Adjusting for confounding with text matching”. In: *American Journal of Political Science* 4 (2020).

[Sah+19] K. Saha, B. Sugar, J. Torous, B. Abrahao, E. Kıcıman, and M. De Choudhury. “A social media study on the effects of psychiatric medication use”. In: *Proceedings of the International AAAI Conference on Web and Social Media*. 2019.

[San+19] V. Sanh, L. Debut, J. Chaumond, and T. Wolf. “Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter”. In: *arXiv preprint arXiv:1910.01108* (2019).

[SBV19] C. Shi, D. M. Blei, and V. Veitch. “Adapting neural networks for the estimation of treatment effects”. In: *Advances in Neural Information Processing Systems*. 2019.

[SG19] D. Sridhar and L. Getoor. “Estimating causal effects of tone in online debates”. In: *arXiv preprint arXiv:1906.04177* (2019).

[Sri+18] D. Sridhar, A. Springer, V. Hollis, S. Whittaker, and L. Getoor. “Estimating causal effects of exercise from mood logging data”. In: *IJCAI/ICML Workshop on CausalML*. 2018.

[Tab+18] N. Tabari, P. Biswas, B. Praneeth, A. Seyeditabari, M. Hadzikadic, and W. Zadrozny. “Causality analysis of twitter sentiments and stock market returns”. In: *Proceedings of the first workshop on economics and natural language processing*. 2018.

[TLP14] C. Tan, L. Lee, and B. Pang. “The effect of wording on message propagation: topic-and author-controlled natural experiments on twitter”. In: *arXiv preprint arXiv:1405.1438* (2014).

[VSB19] V. Veitch, D. Sridhar, and D. M. Blei. “Using text embeddings for causal inference”. In: *arXiv e-prints* (2019).

[WC19] Z. Wang and A. Culotta. “When do words matter? understanding the impact of lexical choice on audience perception using individual treatment effect estimation”. In: *Proceedings of the AAAI Conference on Artificial Intelligence*. 01. 2019.

[WDSD18] Z. Wood-Doughty, I. Shpitser, and M. Dredze. “Challenges of using text classifiers for causal inference”. In: *Proceedings of the Conference on Empirical Methods in Natural Language Processing. Conference on Empirical Methods in Natural Language Processing*. NIH Public Access. 2018.

[YKT18] M. Yeomans, A. Kantor, and D. Tingley. “The politeness package: detecting politeness in natural language.” In: *R Journal* 2 (2018).

[ZMDNM20] J. Zhang, S. Mullainathan, and C. Danescu-Niculescu-Mizil. “Quantifying the causal effects of conversational tendencies”. In: *Proceedings of the ACM on Human-Computer Interaction* CSCW2 (2020).

A Proof of Asymptotic Normality

Theorem 2. *Assume the following.*

1. *The mis-estimation of conditional outcomes can be bounded as follows*

$$\max_{a\in\{0,1\}} \mathbb{E}[(\hat{Q}_a(X)-Q(a,X))^2]^{\frac{1}{2}} = o(n^{-\frac{1}{4}}). \tag{4.9}$$

2. *The propensity score function $P(A=1|\cdot,\cdot)$ is Lipschitz continuous on $\mathbb{R}^2$, and $\exists\ \varepsilon > 0$, $\mathrm{P}\big(\varepsilon \le g_\eta(X) \le 1-\varepsilon\big) = 1$*
3. *The propensity score estimate converges at least as quickly as k nearest neighbor; i.e., $\mathbb{E}[\big(\hat{g}_\eta(X) - \mathrm{P}(A=1 \mid \hat{\eta}(X)\big)^2 \mid X]^{\frac{1}{2}} = O(n^{-\frac{1}{4}})$ [Gyö+02];*
4. *There exist positive constants C_1, C_2, c, and $q > 2$ such that*

$$\begin{aligned}
&\mathbb{E}[|Y|^q]^{\frac{1}{q}} \le C_2,\\
&\sup_{\eta\in supp(\eta(X))} \mathbb{E}[(Y-Q(A,X)^2 \mid \eta(X)=\eta)] \le C_2,\\
&\mathbb{E}[(Y-Q(A,X)^2)]^{\frac{1}{2}} \ge c,\\
&\max_{a\in\{0,1\}} \mathbb{E}[\big|\hat{Q}_a(X)-Q(a,X)\big|]^{\frac{1}{q}} \le C_1.
\end{aligned}$$

Then, the estimator $\hat{\tau}^{TI}$ is consistent and

$$\sqrt{n}(\hat{\tau}^{TI} - \tau^{\mathrm{NDE}}) \xrightarrow{d} \mathbb{N}(0,\sigma^2) \tag{4.10}$$

where $\sigma^2 = E\left(\phi(X;Q,g_\eta,\tau^{\mathrm{NDE}})\right)^2$.

Proof. We first prove that misestimation of propensity score has rate $n^{-\frac{1}{4}}$. For simplicity, we use f_g, $\hat{f}_g\colon (u,v)\in\mathbb{R}^2 \to \mathbb{R}$ to denote conditional probability $\mathrm{P}(A=1|u,v) = f_g(u,v)$ and the estimated propensity function by running the nonparametric regression $\hat{\mathrm{P}}(A=1|u,v) = \hat{f}_g(u,v)$. Specifically, we have $f_g(Q(0,X),Q(1,X)) = g_\eta(X)$ and $\hat{f}_g(\hat{Q}_0(X),\hat{Q}_1(X)) = \hat{\mathrm{P}}(A=1|\hat{Q}_0(X),\hat{Q}_1(X)) = \hat{g}_\eta(X)$. Since $\mathbb{E}[\big(\hat{Q}_0(X)-Q(0,X)\big)^2]^{\frac{1}{2}}$, $\mathbb{E}[\big(\hat{Q}_1(X)-Q(1,X)\big)^2]^{\frac{1}{2}} = o(n^{-1/4})$ and f_g is Lipschitz continuous, we have

$$\begin{aligned}
&\mathbb{E}\left[\left|f_g(\hat{Q}_0(X),\hat{Q}_1(X)) - f_g\left(Q(0,X),Q(1,X)\right)\right|^2\right]^{\frac{1}{2}}\\
\le& L\cdot\mathbb{E}\left[\left\|\big(\hat{Q}_0(X),\hat{Q}_1(X)\big) - \big(Q(0,X),Q(1,X)\big)\right\|_2^2\right]^{\frac{1}{2}}\\
=& L\cdot\left\{\mathbb{E}\left[\big(\hat{Q}_0(X)-Q(0,X)\big)^2\right] + \mathbb{E}\left[\big(\hat{Q}_1(X)-Q(1,X)\big)^2\right]\right\}^{\frac{1}{2}}\\
=& o(n^{-1/4})
\end{aligned} \tag{A.1}$$

Since the true propensity function f_g is Lipschitz continuous on $\mathbb{R}^2$, the mean squared error rate of the k nearest neighbor is $O(n^{-1/2})$ [Gyö+02]. In addition, since the propensity score function and its estimation are bounded under 1, we have the following equation

$$\mathbb{E}\left|\hat{f}_g(\hat{Q}_0(X),\hat{Q}_1(X)) - f_g(\hat{Q}_0(X),\hat{Q}_1(X))\right|^2 = O(n^{-1/2}), \tag{A.2}$$

due to the dominated convergence theorem. By (A.1) and (A.2), we can bound the mean squared error of estimated propensity score in the following form:

$$\begin{aligned}
&\mathbb{E}\left[\left(\hat{g}_\eta(X)-g_\eta(X)\right)^2\right] \\
\leq&\mathbb{E}\left[\left(\hat{g}_\eta(X)-f_g(\hat{Q}_0(X),\hat{Q}_1(X))\right)^2\right]+\mathbb{E}\left[\left(f_g(\hat{Q}_0(X),\hat{Q}_1(X))-g_\eta(X)\right)^2\right] \\
=&\mathbb{E}\left|f_g(\hat{Q}_0(X),\hat{Q}_1(X))-f_g\left(Q(0,X),Q(1,X)\right)\right|^2+ \\
&\mathbb{E}\left|\hat{f}_g(\hat{Q}_0(X),\hat{Q}_1(X))-f_g(\hat{Q}_0(X),\hat{Q}_1(X))\right|^2 \\
=&O(n^{-1/2}),
\end{aligned} \tag{A.3}$$

that is $\mathbb{E}\left[\left(\hat{g}_\eta(X)-g_\eta(X)\right)^2\right]=O(n^{-\frac{1}{4}})$.

Before we apply the conclusion of Theorem 5.1 in [Che+17b], we need to check all assumptions in Assumption 5.1 hold in Chernozhukov et al. [Che+17b]. Let $C:=\max\left\{(2C_1^q+2^q)^{\frac{1}{q}},C_2\right\}$.

(a) $\mathbb{E}[Y-Q(A,X)\mid\eta(X),A]=0$, $\mathbb{E}[A-g_\eta(X)\mid\eta(X)]=0$ are easily checked by invoking definitions of Q and g_η.

(b) $\mathbb{E}[|Y|^q]^{\frac{1}{q}}\leq C$, $\mathbb{E}[(Y-Q(A,X))^2]^{\frac{1}{2}}\geq c$ and $\sup_{\eta\in\mathrm{supp}(\eta(X))}\mathbb{E}[(Y-Q(A,X))^2\mid\eta(X)=\eta]\leq C$ are guaranteed by the fourth condition in the theorem.

(c) $\mathrm{P}\left(\varepsilon\leq g_\eta(X)\leq 1-\varepsilon\right)=1$ is the second condition in the theorem.

(d) Since propensity score function and its estimation are bounded under 1, we have

$$\begin{aligned}
&\left(\mathbb{E}[\left|\hat{Q}_1(X)-Q(1,X)\right|^q]+\mathbb{E}[\left|\hat{Q}_0(X)-Q(0,X)\right|^q]+\mathbb{E}[\left|\hat{g}_\eta(X)-g_\eta(X)\right|^q]\right)^{\frac{1}{q}} \\
&\leq(C_1^q+C_1^q+2^q)^{\frac{1}{q}} \\
&\leq C
\end{aligned}$$

(e) Based on (A.3) and condition 1 in the theorem, we have

$$\begin{aligned}
&\left(\mathbb{E}[\left(\hat{Q}_1(X)-Q(1,X)\right)^2]+\mathbb{E}[\left(\hat{Q}_0(X)-Q(0,X)\right)^2]+\mathbb{E}[\left(\hat{g}_\eta(X)-g_\eta(X)\right)^2]\right)^{\frac{1}{2}} \\
&\leq\left[o(n^{-\frac{1}{2}})+o(n^{-\frac{1}{2}})+O(n^{-\frac{1}{2}})\right]^{\frac{1}{2}} \\
&\leq O(n^{-\frac{1}{4}}), \\
&\mathbb{E}[\left(\hat{Q}_0(X)-Q(0,X)\right)^2]^{\frac{1}{2}}\cdot\mathbb{E}[\left(\hat{g}_\eta(X)-g_\eta(X)\right)^2]^{\frac{1}{2}}=o(n^{-\frac{1}{2}})
\end{aligned}$$

(f) Based on condition 3 in the theorem, we have

$$\sup_{x\in\mathrm{supp}(X)}\mathbb{E}[\left(\hat{g}_\eta(X)-\mathrm{P}(A=1\mid\hat{\eta}(X))^2\mid X=x\right)^2]=O(n^{-\frac{1}{2}}),$$

where $\mathrm{supp}(X)$ is the support of the random variable X. We consider a smaller positive

constant $\tilde{\varepsilon}$ instead of ε. Note that for $\tilde{\varepsilon} < \varepsilon$, we still have $\mathrm{P}(\tilde{\varepsilon} \le g_\eta(X) \le 1-\tilde{\varepsilon}) = 1$. Then,

$$
\begin{aligned}
&\mathrm{P}\left(\sup_{x\in\mathrm{supp}(X)}\left|\hat{g}_\eta(x)-\frac{1}{2}\right| > \frac{1}{2}-\tilde{\varepsilon}\right) = \mathrm{P}\left(\inf_{x\in\mathrm{supp}(X)}\hat{g}_\eta(x) < \tilde{\varepsilon}\right) + \mathrm{P}\left(\sup_{x\in\mathrm{supp}(X)}\hat{g}_\eta(x) > 1-\tilde{\varepsilon}\right)\\
\le&\mathrm{P}\left(\inf_{x\in\mathrm{supp}(X)}\mathrm{P}(A=1\mid\hat{\eta}(X)=\eta(x)) - \inf_{x\in\mathrm{supp}(X)}\hat{g}_\eta(x) > \varepsilon-\tilde{\varepsilon}\right)\\
&+\mathrm{P}\left(\sup_{x\in\mathrm{supp}(X)}\hat{g}_\eta(x) - \sup_{x\in\mathrm{supp}(X)}\mathrm{P}(A=1\mid\hat{\eta}(X)=\eta(x)) > 1-\tilde{\varepsilon}-(1-\varepsilon)\right)\\
\le&\frac{\mathbb{E}\left[\left(\inf_{x\in\mathrm{supp}(X)}\hat{g}_\eta(x)-\inf_{x\in\mathrm{supp}(X)}\mathrm{P}(A=1\mid\hat{\eta}(X)=\eta(x))\right)^2\right]}{(\varepsilon-\tilde{\varepsilon})^2}+\\
&\frac{\mathbb{E}\left[\left(\sup_{x\in\mathrm{supp}(X)}\hat{g}_\eta(x)-\sup_{x\in\mathrm{supp}(X)}\mathrm{P}(A=1\mid\hat{\eta}(X)=\eta(x))\right)^2\right]}{(\varepsilon-\tilde{\varepsilon})^2}\\
\le&\frac{2\sup_{x\in\mathrm{supp}(X)}\mathbb{E}\left[\left(\hat{g}_\eta(X)-\mathrm{P}(A=1\mid\hat{\eta}(X))^2\mid X=x\right)^2\right]}{(\varepsilon-\tilde{\varepsilon})^2}\\
=&O(n^{-\frac{1}{2}})
\end{aligned}
$$

Hence, $\mathrm{P}(\sup_{x\in\mathrm{supp}(X)}\left|\hat{g}_\eta(x)-\frac{1}{2}\right| \le \frac{1}{2}-\tilde{\varepsilon}) \ge 1-O(n^{-\frac{1}{2}})$.

With (a)-(f), we can invoke the conclusion in Theorem 5.1 in [Che+17b], and get the asymptotic normality of the TI estimator. □

B Proof of Causal Identification

Theorem 1. *Assume the following:*

1. *(Causal structure) The causal relationships among A, $\tilde{A}$, Z, Y, and X satisfy the causal DAG in Figure 2*
2. *(Overlap)* $0 < \mathrm{P}(A=1|X_{A\wedge Z}, X_Z) < 1$
3. *(Intention equals perception)* $A = \tilde{A}$ *almost surely.*

Then, the NDE is identified from observational data as

$$\mathrm{NDE} = \tau^{\mathrm{NDE}} := \mathbb{E}_{X|\tilde{A}=1}\left[\mathbb{E}[Y\mid\eta(X),\tilde{A}=1] - \mathbb{E}[Y\mid\eta(X),\tilde{A}=0]\right], \tag{3.4}$$

where $\eta(X) := (f(Q(0,X)), h(Q(1,X)))$*, and* $f, h: \mathbb{R}\to\mathbb{R}$ *are any invertible functions on* $\mathbb{R}$*.*

Proof. Since f and h are two invertible functions on $\mathbb{R}$, the sigma algebra should be the same for $(Q(0,X), Q(1,X))$ and $f(Q(0,X)), h(Q(1,X))$, i.e.,

$$\sigma(Q(0,X), Q(1,X)) = \sigma(f(Q(0,X)), h(Q(1,X))).$$

Hence, we have

$$
\begin{aligned}
&\mathrm{P}(A=1|\,Q(0,X), Q(1,X)) = \mathrm{P}(A=1|\,f(Q(0,X)), h(Q(1,X))),\\
&\mathbb{E}(Y\mid\,Q(0,X), Q(1,X)) = \mathbb{E}(Y\mid\,f(Q(0,X)), h(Q(1,X))),
\end{aligned}
\tag{B.1}
$$

and we only need to prove $\eta(X) = (Q(0,X), Q(1,X))$ suffices the identification for NDE.

We first prove that this two-dimensional confounding part $\eta(X)$ satisfies overlap condition. Since $(Q(0,X),\ Q(1,X)) = (\mathbb{E}[Y\mid A=1, X_{A\wedge Z}, X_Z],\ \mathbb{E}[Y\mid A=0, X_{A\wedge Z}, X_Z])$ is a function of

$(X_{A\wedge Z},\ X_Z)$, the following equations hold:

$$\begin{aligned}\mathrm{P}(A=1 \mid Q(0,X),Q(1,X)) =& \mathbb{E}(A \mid Q(0,X),Q(1,X)) \\ =& \mathbb{E}\left[E\left(A \mid X_{A\wedge Z},X_Z\right) \mid Q(0,X),Q(1,X)\right] \\ =& \mathbb{E}\left[\mathrm{P}(A=1 \mid X_{A\wedge Z},X_Z) \mid Q(0,X),Q(1,X)\right].\end{aligned} \tag{B.2}$$

As $0<\mathrm{P}(A=1 \mid X_{A\wedge Z},X_Z)<1$, we have $0<\mathrm{P}(A=1 \mid Q(0,X),Q(1,X))<1$. Furthermore, we have $0<\mathrm{P}(\tilde{A}=1 \mid Q(0,X),Q(1,X))<1$ due to almost everywhere equivalence of A and $\tilde{A}$.

Since $A=\tilde{A}$, we can rewrite (3.1) by replacing A with $\tilde{A}$ in the following form:

$$\begin{aligned}\mathrm{NDE} =& \mathbb{E}_{X_{A\wedge Z},X_Z \mid \tilde{A}=1}\left[\ \mathbb{E}(Y \mid \mathrm{do}(\tilde{A}=1),X_{A\wedge Z},X_Z)-\mathbb{E}(Y \mid \mathrm{do}(\tilde{A}=0),X_{A\wedge Z},X_Z)\ \right] \\ =& \mathbb{E}_{X_{A\wedge Z},X_Z \mid \tilde{A}=1}\left[\ \mathbb{E}(Y \mid \tilde{A}=1,X_{A\wedge Z},X_Z)-\mathbb{E}(Y \mid \tilde{A}=0,X_{A\wedge Z},X_Z)\ \right] \\ =& \mathbb{E}_{X_{A\wedge Z},X_Z \mid \tilde{A}=1}\left[\ \mathbb{E}(Y \mid \tilde{A}=1,X)-\mathbb{E}(Y \mid \tilde{A}=0,X)\ \right] \\ =& \mathbb{E}_{X_{A\wedge Z},X_Z \mid \tilde{A}=1}\left[\mathbb{E}(Y \mid \tilde{A}=1,Q(0,X),Q(1,X))\right]-\mathbb{E}\left[\mathbb{E}(Y \mid \tilde{A}=0,Q(0,X),Q(1,X))\right] \\ =& \mathbb{E}_{X_{A\wedge Z},X_Z \mid \tilde{A}=1}\left[\mathbb{E}(Y \mid \tilde{A}=1,\eta(X))\right]-\mathbb{E}\left[\mathbb{E}(Y \mid \tilde{A}=0,\eta(X))\right].\end{aligned} \tag{B.3}$$

The equivalence of the first and the second line is because $X_{A\wedge Z}$, X_Z block all backdoor paths between $\tilde{A}$ and Y (See Figure 2) and $0<\mathrm{P}(\tilde{A}=1 \mid Q(0,X),Q(1,X))<1$. Thus, the "*do-operation*" in the first line can be safely removed. Equivalence of the second line and the third line is due to $Q(\tilde{A},X)=\mathbb{E}\left(Y \mid \tilde{A},X_{A\wedge Z},X_Z\right)$, which is subject to the causal model in Figure 2.

(B.3) shows that $(Q(0,X),Q(1,X))$ is a two-dimensional confounding variable such that NDE is identifiable when we adjust for it as the confounding part.

□

C Additional Experiments

We conduct additional experiments to show how the estimation of causal effect changes 1) over different nonparametric models for the propensity score estimation, and 2) when using different double machine learning estimators on causal estimation. Specifically, for the first study, we apply different nonparametric models and the logistic regression to the estimated confounding part $\hat{\eta}(X)=\left(\hat{Q}_0(X),\hat{Q}_1(X)\right)$ to obtain propensity scores. We use ATT AIPTW in all above cases for causal effect estimation. For the second study, we fix the first two stages of the TI estimator, i.e. we apply Q-Net for the conditional outcomes and compute propensity scores with the Gaussian process regression where the kernel function is the summation of dot product and white noise. Estimated conditional outcomes and propensity scores are plugged into different double machine learning estimators. We make the following conclusions with results of above experiments.

The choice of nonparametric models is significant. Table 3 summarizes results with applying different regression models for the propensity estimation. We can see that suitable nonparametric models will strongly increase the coverage proportion over true causal estimand. Therefore, we conclude that the accuracy in causal estimation is highly dependent on the choice of nonparametric models. In practice, when there is some prior information about the propensity score function, we should apply the most suitable nonparametric model to increase the reliability of our causal estimation.

The ATT AIPTW is consistently the best double machine learning estimator. Table 4 shows results by applying different double machine learning estimators. We apply both estimators for the average treatment effect (ATE) and the natural direct effect (NDE). The bias of "unadjusted" estimator $\hat{\tau}^{\mathrm{naive}}$ and $\hat{\tau}^{\mathrm{naive+C}}$ are also included in Table 4 (a). For bias, ATT AIPTW $\hat{\tau}^{\mathrm{TI}}$ has comparable results with other double machine learning estimators in most cases. For

coverage proportion of confidence intervals, though it has lower rates in some cases, $\hat{\tau}^{\mathrm{Q}}$ has consistently the best performance. Especially in high confounding situations, the advantage of $\hat{\tau}^{\mathrm{Q}}$ is obvious.

Estimator For each dataset, we compute estimators as follows. n_1 and n_0 stands for the number of individuals in the treated and controlled group. $n = n_1 + n_0$ is the total number of individuals.

- "Unadjusted" baseline estimator: $\hat{\tau}^{\mathrm{naive}} = \frac{1}{n_1}\sum_{i:A_i=1} Y_i - \frac{1}{n_0}\sum_{i:A_i=0} Y_i$
- "Outcome-only" estimator: $\hat{\tau}^{\mathrm{Q}} = \frac{1}{n_1}\sum_{i:A_i=1} \hat{Q}_{1,i} - \hat{Q}_{0,i}$
- ATT AIPTW: $\hat{\tau}^{\mathrm{TI}} = \frac{1}{n_1}\sum_{i:A_i=1} A_i(Y_i - \hat{Q}_{0,i}) - (1-A_i)(Y_i - \hat{Q}_{0,i})\frac{g_i}{1-\hat{g}_i}$

Table 3: The choice of nonparametric models for the TI-estimator is significant. Tables show average bias and 95% confidence intervals' coverage of $\hat{\tau}^{\mathrm{TI}}$ with applying different nonparametric models in the second stage. The Gaussian process regression with the dot product+ white noise kernel has the best performance (lowest bias and highest coverage proportion). The treatment level is equal to true NDE, which takes 1.0 (with causal effect) and 0.0 (without causal effect). Low and high noise level corresponds to $\gamma = 1.0$ and 4.0. Low and high confounding level corresponds to $\beta_c = 50.0$ and 100.0.

(a) Average bias

Noise:	Low				High			
Treatment (oracle causal effect):	1.0		0.0		1.0		0.0	
Confounding:	Low	High	Low	High	Low	High	Low	High
GPR (Dot Product+White Noise)	0.065	0.052	0.104	0.070	0.077	0.044	-0.038	0.084
GPR (RBF)	0.148	0.346	0.150	0.326	0.353	0.440	0.332	0.412
KNN	0.144	0.331	0.140	0.310	0.308	0.364	0.297	0.352
AdaBoost	0.067	0.341	0.055	0.319	0.512	0.477	0.466	0.450
Logistic	0.065	0.050	0.105	0.068	0.075	0.043	-0.042	0.083

(b) Coverage proportions of 95% confidence intervals

Noise:	Low				High			
Treatment (oracle causal effect):	1.0		0.0		1.0		0.0	
Confounding:	Low	High	Low	High	Low	High	Low	High
GPR (Dot Product+White Noise)	60%	86%	60%	86%	86%	88%	80%	84%
GPR (RBF)	34%	0%	42%	0%	10%	10%	24%	26%
KNN	26%	0%	40%	0%	16%	12%	14%	12%
AdaBoost	34%	0%	46%	0%	0%	0%	0%	0%
Logistic	60%	86%	60%	86%	86%	88%	80%	84%

Table 4: The ATT AIPTW is consistently the best double machine learning estimator for this causal problem. Tables show average bias and 95% confidence intervals' coverage of different causal estimations. ATT AIPTW $\hat{\tau}^{\mathrm{TI}}$ shows consistently the lowest bias and highest coverage rate. For propensity score estimation, the Gaussian process regression with the dot product+ white noise kernel is applied for all estimators. The treatment level is equal to true NDE/true ATE, which takes 1.0 (with causal effect) and 0.0 (without causal effect). Low and high noise level corresponds to $\gamma = 1.0$ and 4.0. Low and high confounding level corresponds to $\beta_c = 50.0$ and 100.0.

(a) Average bias

Noise:	Low				High			
Treatment (oracle NDE):	1.0		0.0		1.0		0.0	
Confounding:	Low	High	Low	High	Low	High	Low	High
unadjusted $\hat{\tau}^{naive}$	1.073	2.145	1.073	2.145	1.079	2.150	1.079	2.150
ATE AIPTW	0.090	0.171	0.110	0.196	0.112	0.103	0.031	0.132
ATE BMM	0.089	0.169	0.110	0.194	0.112	0.102	0.031	0.131
ATE IPTW	-0.645	-1.577	-1.964	-1.837	-0.117	-0.646	-0.504	-0.678
ATT AIPTW: $\hat{\tau}^{TI}$	0.065	0.052	0.104	0.070	0.077	0.044	-0.038	0.084
ATT BMM	0.065	0.137	-0.036	0.057	0.614	0.435	0.457	0.329

(b) Coverage Proportions of 95% confidence intervals

Noise:	Low				High			
Treatment (oracle NDE):	1.0		0.0		1.0		0.0	
Confounding:	Low	High	Low	High	Low	High	Low	High
ATE AIPTW	38%	40%	74%	32%	78%	84%	82%	72%
ATE BMM	42%	40%	78%	34%	80%	84%	82%	72%
ATE IPTW	14%	2%	0%	2%	88%	34%	32%	36%
ATT AIPTW: $\hat{\tau}^{TI}$	60%	86%	60%	86%	86%	88%	80%	84%
ATT BMM	38%	8%	38%	40%	2%	6%	2%	18%